

Robust Charging Infrastructure Design under Demand Uncertainty

A solar-powered charging station for a shared e-micromobility fleet — Executive Summary

Competition Themes — 3 (primary): solar-PV charging-station design · 2 (secondary): charge/discharge profiles under demand scenarios

20 kWp PV · 0 kWh · 8 chargers Recommended robust design	-65% (£91,499/yr) Value of robustness (worst-case cost)	97% (vs 86% naive) Guaranteed fleet service, worst case	16% (4.6 tCO2/yr) Operational carbon saving
--	---	---	---

1. Problem Understanding [20% criterion]

A shared-micromobility operator depot charges a mixed fleet of e-scooters, e-bikes and e-cargo bikes whose demand swings with season, weather, events and multi-year growth. This creates a sizing dilemma for a solar charging hub: one built for **average** demand fails at peaks and strands the fleet, while one built for the **worst case** sinks capital into rarely-used equipment. The dilemma sharpens under a **constrained grid connection** — a higher import limit needs a costly DNO reinforcement — so the evening charging peak must be served from on-site solar and storage, across enough charge bays. Sizing becomes a decision under deep uncertainty: under-build and lose service; over-build and waste money. Demand is anchored to the REAL DfT Newcastle e-scooter trial (monthly trips, fleet size, deployment, trip distance), so the problem is posed on observed data, not assumption.

2. Relevance & Innovativeness of the Proposed Solution [20% criterion]

The solution directly delivers **Theme 3** (design of a solar-PV charging station for a shared fleet) and **Theme 2** (modelling charge/discharge profiles under demand scenarios), producing an actionable station specification with uncertainty bounds. Its novelty is methodological: charging-station sizing is almost always **deterministic** (a single demand forecast). We instead apply **robust optimisation — minimax-regret, maximin and CVaR — alongside a two-stage stochastic program**, choosing a design that performs across the entire demand space, and we define and quantify the “**value of robustness**”: the worst-case cost and lost service avoided by hedging. A correlated Monte-Carlo “demand regime” reproduces the realistic co-movement of trips, utilisation and growth, and exposes a subtle but important result — expected-value design **under-weights the correlated high-demand tail**, which is precisely why scenario-based robust methods are required here. A further, counter-intuitive finding falls out: because depot charging is flexible, **smart charging is the cheapest robustness lever** — it flattens the load below the connection so that, at a grid-connected depot, a battery does not pay; storage becomes essential only as the connection weakens toward off-grid (~10 kW), a boundary we quantify. The recommendation is therefore solar + smart-managed bays, not capital-heavy storage.

3. Approach — Methodology, Algorithm & Flow [20% criterion]

The pipeline (Figure 1) is: **(i)** build Low/Medium/High demand scenarios from the real trial data, scaled across a year-on-year growth range into nine probability-weighted scenarios; **(ii)** simulate an 8,760-hour energy balance for any design under the capped grid, dispatching PV → battery (with off-peak grid pre-charging for peak-shaving) → capped grid → unmet demand; **(iii)** evaluate all 150 designs (PV 5–25 kWp × battery 0–50 kWh × 4–20 chargers) against the nine scenarios, pricing unmet demand as lost service, to form a cost matrix; **(iv)** apply five decision rules — naive, a two-stage stochastic program (expected cost), CVaR (risk-averse), minimax-regret and maximin — and map the cost-vs-robustness Pareto frontier; **(v)** propagate uncertainty through a 500-sample correlated Monte-Carlo fan and rank drivers with a tornado and variance-based global (Sobol) sensitivity; **(vi)** confirm the conclusion is not an artefact via a robustness-of-robustness sweep (it holds across every penalty × grid combination) and validate seven outputs against published benchmarks. The full study runs in under a minute.

Methodology — Robust Solar Charging-Station Design




Figure 1 — End-to-end methodology / algorithm flow.

4. Feasibility of the Proposed Approach [20% criterion]

Technical: the design uses only commercially available components — monocrystalline PV, a Li-ion battery (0–50 kWh) and standard AC charge points — sized within ranges supported by manufacturer datasheets and industry benchmarks. **Economic:** the recommended robust station has a ≈£36,000 capital cost and a bounded annual cost even in the worst demand scenario; the model reports CAPEX, OPEX, battery-replacement and grid costs. **Operational:** it yields one deployable specification plus the confidence interval around it. **Validated:** seven key outputs (PV yield, carbon intensity, round-trip efficiency, solar fraction, LCOE, carbon saving, trip distance) all fall within published ranges. **Computational:** the model is open-source Python, deterministic, regenerates all data and figures from one command, and passes an automated test suite — any reviewer can reproduce every number.

5. Originality & Authenticity (AI / Plagiarism) [20% criterion]

All model code, the dispatch engine, the optimisation formulation and every figure were written from scratch for this submission. All three datasets are REAL: DfT e-scooter monitoring data (Open Government Licence), hourly solar pulled live from the EU PVGIS API, and regional grid-carbon intensity from the National Grid ESO API. Every numerical assumption carries an inline source citation and all literature is referenced. The repository is fully self-contained and reproducible — fixed random seed, one-command rebuild, passing tests — directly supporting the competition's AI and plagiarism checks. No third-party text or code is reproduced.





Figure 2 — Cost-vs-robustness Pareto frontier (five decision rules). Figure 3 — Robustness of the conclusion: robust beats naive across every penalty x grid combination.